

Table of content

1. [Unit 1: Overview](#)
2. [Unit 2: ER Model](#)
3. [Unit 3: The Relational Model and Relational Database Design](#)

Unit 1

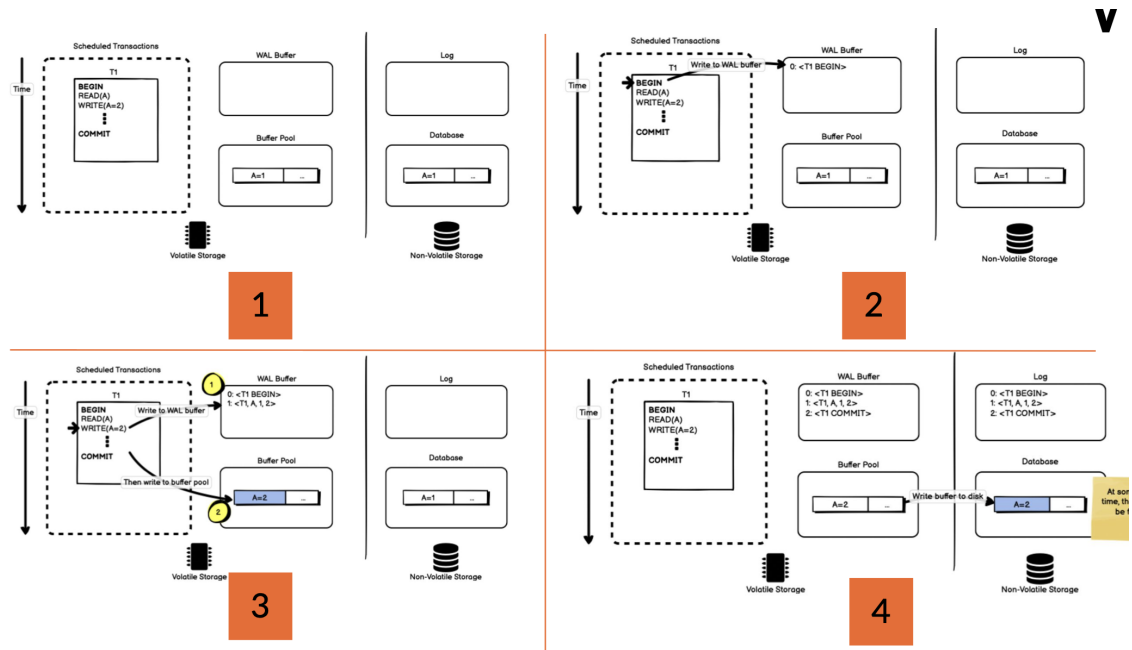
Overview

Terminology :

- **Data** - 50 cats
- **Database**
 - collection of data which can accessed from computer system
 - The data in the same database should relate have "entities" along with "relationships"
- **Database machines** - or back end processer is the hardware for database
- **DBMS** - Database Management - software package for database

DBMS Character

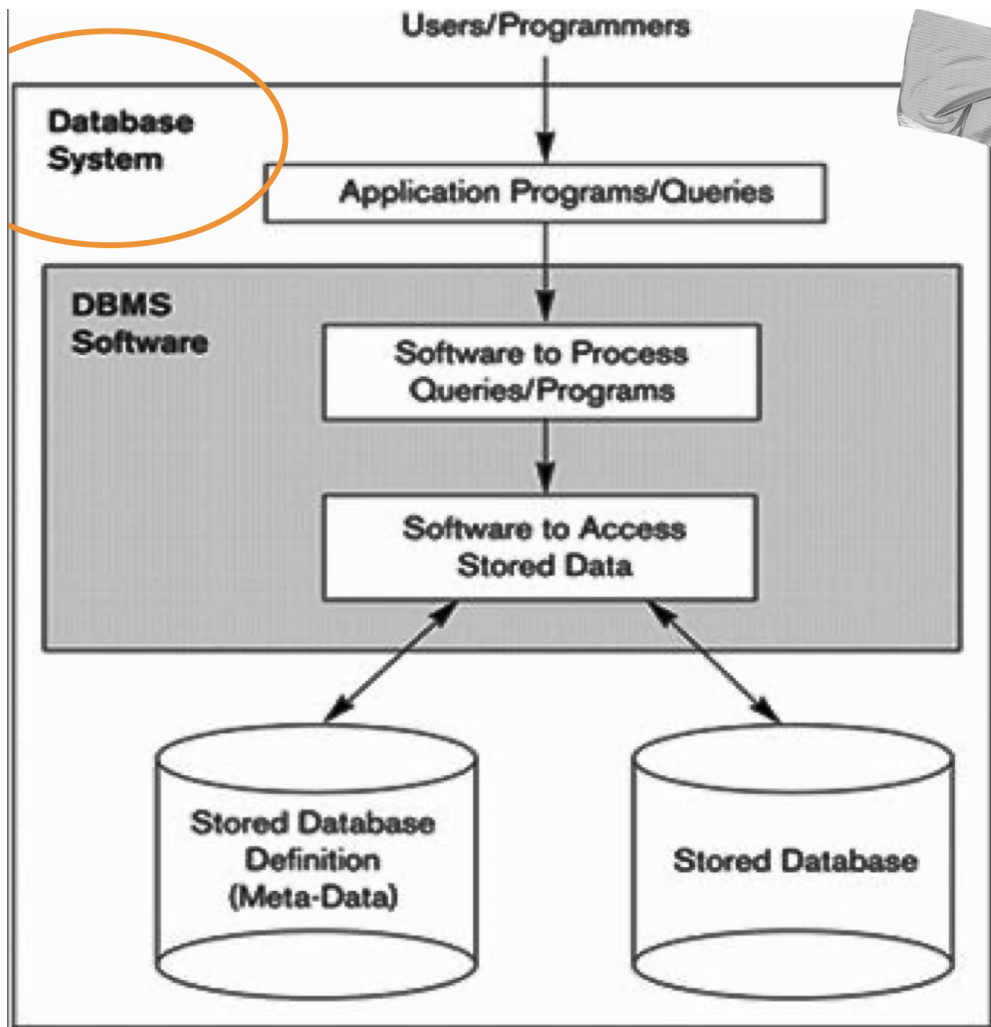
- **less redundancy** - the same piece of data store many place
- **data consistency** - accurate and up-to-date across difference system
- **data integrity** - unsure that data remain intact, uncorrupt
- **data security** - implement strict access rule, encrytion
- **multi-user support** - concurrency control
- **backup & recovery**
- **query language support**
- **ACID transaction** (not all DBMS can do all the ACID)
 - **A: Atomicity** - transactions are all or nothing
 - **C: Consistency** - Only valid data will be saved
 - **I: Isolation** - transactions don't affect each others
 - **D: Durability** - written data will not be lost
 - How are ACID transactions implemented
 - most via "lock"
 - to guarantee duability use "write-ahead lock"
 - rolled back or continued from the transaction log left off



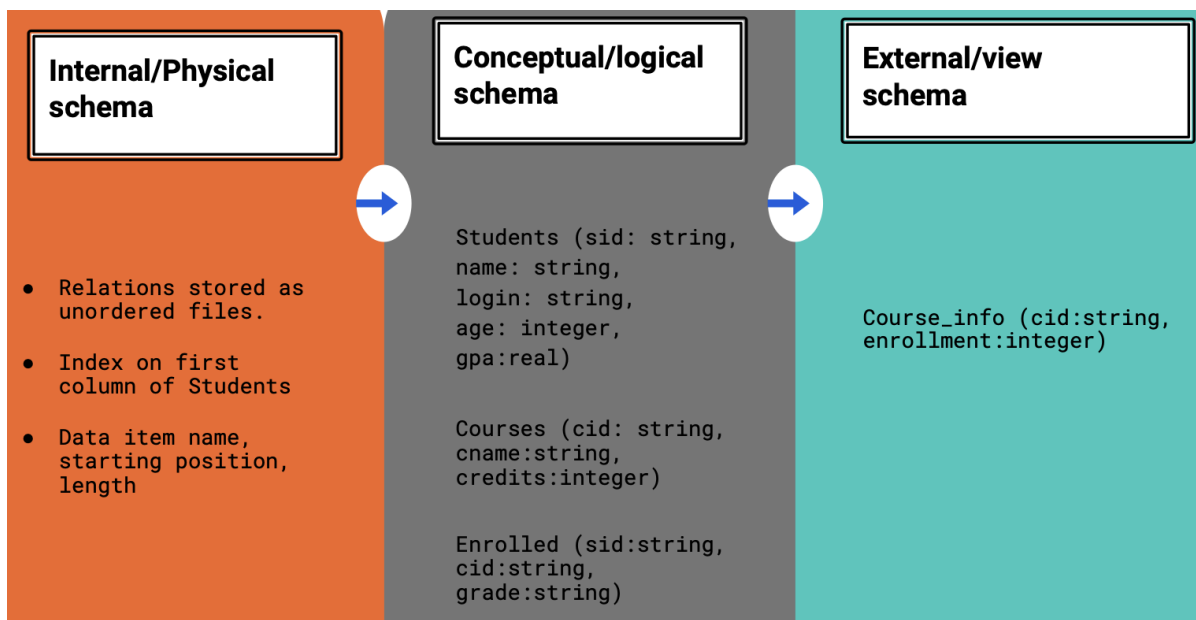
CAP theorem or Brewer's theorem

- impossible for a dist data store to simultaneously provide more than two of
 - Consistency:** when system return info, it is always up-to-date
 - Availability:** systems always return info, even if state
 - Partition tolerance:** system continues operating during a patition

A Simplified Database System Environment



Level of Abstraction Three Schema Architecture



1. **Internal Level** - low-level data structure
2. **Conceptual Level** - what data to store and what relationships
3. **External Level** - difference view of users

SQL Language

- **DDL** (Data Definition Language)
 - create, drop, comment
- **DML** (Data Manipulation Language)
 - select, update, insert, call
- **DCL** (Data Control Language)
 - grant, revoke
- **TCL** (Transaction Control Language)
 - commit, rollback

NoSQL = Not Only SQL or Non-SQL NoSQL is non-relational database, compatible with Big data

	RDBMS	NoSQL
Type	Relational	Non-Relational
Data Format	Table	JSON(Text) etc.
Scaling	Vertical (increase Spec Server)	Horizontal(increase Server)
Schema	Fixed	Flex
Ex	Oracle, MySQL, Microsoft SQL	MongoDB

Unit 2

ER-model

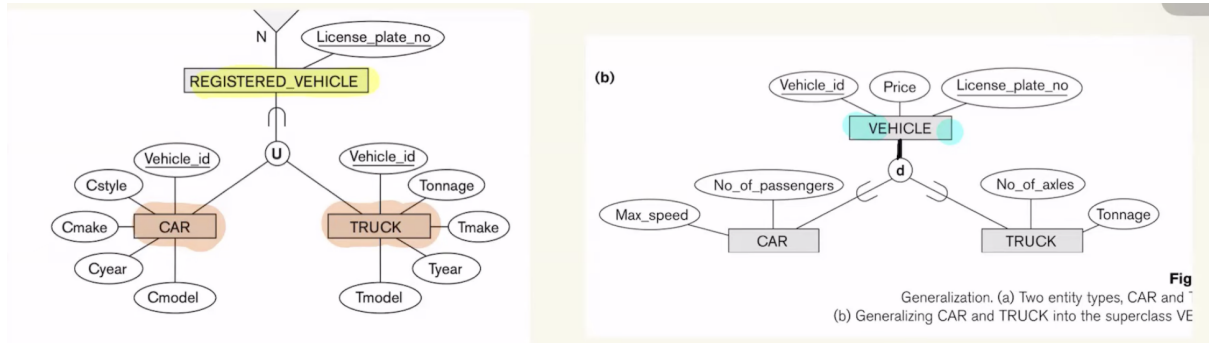
- **Entities** are specific objects in the database
- **Attributes** are properties used to describe any entity
 - Type of Attributes
 1. **Simple** - single atomic ex. SSN, Sex
 2. **Composite** - may be composed of several component ex. Address (Apt#, House#, City)
 3. **Single-valued** - ex. age, height
 4. **Multi-valued** - ex. {car}
 - 1. **Stored attributes** - birth date
 2. **Derived attributes** - age
 3. **Null values** - not have applicable value
- **Entity Types and Key Attributes**
 - ex. *EMPLOYEE* entity type and *SSN of EMPLOYEE* is the key attribute
 - the collection of all entities of a particular entity type at any point in time is called **an entity set**
 - An entity type may have more than one key
 - A key attribute may be composite
- **Relationships and Relationship Types**
 - A relationship relates two or more entities with a specific meaning
 - *EMPLOYEE John Smith works on the ProductX PROJECT*
 - Relationships of the same type group into a relationship type

- EMPLOYEEs WORKS_ON PROJECTs
- **Degree of Relationship Type** - number of participating entity types
- More than one relationship type can exist with the same participating entity types
- can have **recursive** relationship type
 - ex. ManagerID is also EmployeeID
- A relationship type can have attributes
 - ex. work_on -> start_date
- **Constraints on Relationship**
 - Cardinality Ratios (placing number on the link)
 - 1:1, 1:N, N:1, M:N
 - Participation constraint
 - total participation (one or more - mandatory): shown by double lining the link
 - partial participation (zero - optional): shown by single lining the link
 - We will refer to the cardinality ratio and participation constraints, taken together, as the **structural constraints** of a relationship type
- **Weak Entity Types** - entity that **doesn't have key attribute**
 - entity that have a key that matching from other entity
- **Subclasses and Superclasses** - IS-A relationship
 - the same as java
 - use node(d) and line out
 - **Attribute Inheritance**
 - subclass inherits all attribute of superclass, also inherits all relationships
 - **Specialization and Generalization**
 - CAR, TRUCK is a specialization of VEHICLE
 - CAR, TRUCK generalize into VEHICLE
 - **Constraints**
 - READTHID
 - **predicate-defined** -?
 - **attribute-defined** -?
 - user-defined -?
 - 2 other conditions apply
 - **Disjointness Constraint**
 - **disjointed**(d) - an entity can be a member of at most 1 subclass
 - **overlap** (o) - can be more than one subclass
 - **Completeness Constraint**
 - **partial** (single line) - every entity in superclass must be a member of some subclass
 - **total** (double line) - allows an entity not to belong to any subclasses
 - **Insertion and Deletion Rules**
 - **Deleting** an entity from superclass -> auto delete from all subclasses
 - **Inserting** an entity in superclass
 - *partial* don't have to insert in subclass

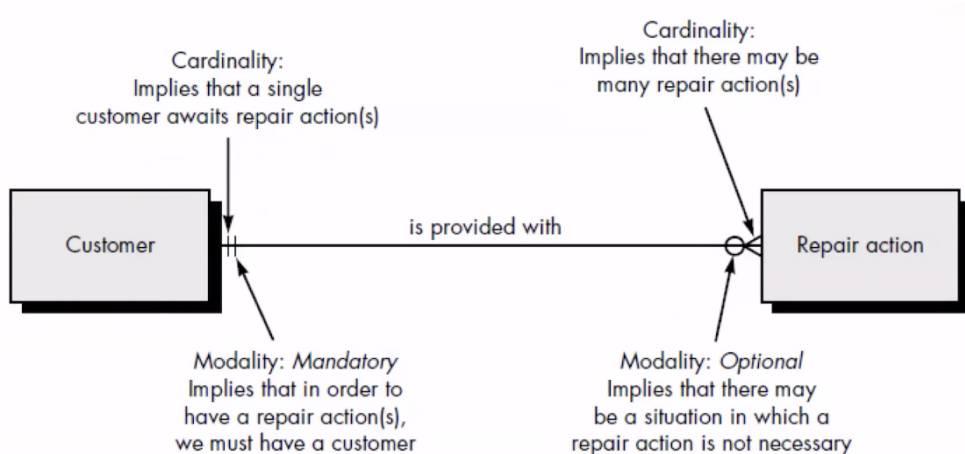
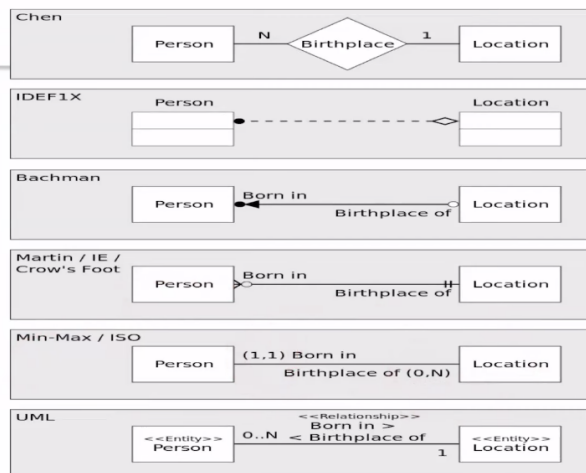
- *total* have to insert into at least one subclass
- **Hierarchies** - only one superclass (single inheritance)
- **Lattice** - can be subclass of more than one superclasses (multiple inheritance)
 - call that subclass "share subclass"

- **Categories (UNION TYPES)**

- a shared subclass is subclass in more than one distinct superclass/subclass relationships, where each relationships has a single superclass



alt="difference between categories and share subclass" ไข่ notation



crow's foot

notation

UNIT 3

Informally, a relation looks like a *table* of values. typically contains a set of rows. Each rows represent certain facts that correspond to a real-world entity or relationship. Each *column* has a header that gives a meaning of the data items.

- **Key of a Relation**

- each row have a *key* (data item or set of items)
- sometimes gen a row-id (not a good choice)

- **Formal Definition**

- $R(A_1, A_2, \dots, A_n) \rightarrow R$ is name of the relation
- Relation schema R defined over attributes A_1, A_2, \dots, A_n
- Each A_i is the name of a role played by some **Domain D** and denoted by $\text{dom}(A_i)$
- the *degree* or *arity* is the number of attributes n of its R
- relation is set of **tuples**(rows), Column called attributes
- $r(R)$: a specific state of R

Informal Terms	Formal Term
table	relation
column header	attribute
all possible column value	domain
row	tuple
table definition	schema of a relation
populated table	state of the relation

- **Chatacteristics of Relation**

- **Ordering of tuples** - the tuple are not considered to be ordered (even they appear to be in the tabular form)
- **Ordering of attributes** - must be order!
- **Values in a tuple**
 - all values are **atomic** (indivisible)
 - a *null* value is used to represent value that are unknow or inapplicable

- **Relational Database Constraints**

- **inherent model-based constraints** - no duplication tuple
- **schema-based constraints** - specify by DDL
 - **Domain constraints**
 - within each tuple, the value of each attribute A must be atomic
 - data types for each attribute
 - a subrange of values, an enumerated type
 - **Key onstraints**
 - all element of a set are distinct
 - denote such the subset by SK (Superkey)
 - SK is a set of attributes that uniquely identifies a row
 - every relation has at least one *default superkey* - the set of all its attribute



- Key of R is a superkey which cannot delete any attribute, to still be superkey
 - if a relation has several **candidate key**, choose primary key (underline) -> the left one is alternate key
 - **Entity integrity onstraints**
 - S is the name of the database -> $S = \{R_1, R_1, \dots, R_n\}$
 - **primary key cannot be null**
 - **Referential integrity onstraints**
 - involving two relation (referencing relation and referenced relation)
 - referencing relation have **FK** (foreign key) that reference the promary key attributes PK of the referenced relation
 - $t_1[FK] = t_2[PK]$
 - **other type**
 - Semantic Integrity Constraints
 - Functional Dependency Constraints
 - **application-based constraints**
- **Update Operation**
 - insert: insert tuple
 - delete: delete tuple
 - update or mondify: change value of existing tuple
 - *integrity constraint should bot be violated by this operation*
 - **insert** - can violate
 - domain constraints: if attribute value violates the corresponding constraint
 - key constraints: dup key value
 - entity integrity: new PK is null
 - referential intergrity: refers to a tuple that not exist
 - **delete** - can violate
 - *only* referential intergrity - the tuple being deleted is referenced by the FK
 - incase that happen
 - reject deletion
 - attempt to cascade(propagate) - delete the tuple that reference
 - modify the referencing attribute values
 - execute a user-specified error-correction routine
 - in some case cannot be null, be careful XD
 - **update** - can violate everything TT

Normal Forms

First Normal Form (1NF)

- **Atomic values only**

Each field should contain only one value.

-  `course_name = "Astrology, Tarot"`
-  `course_name = "Astrology"`

- **No repeating groups**

Avoid multiple similar columns for the same attribute (e.g., `phone1`, `phone2`). Use separate rows or

related tables.

- **Unique rows with a Primary Key**

Each table must have a primary key to uniquely identify each record.

Second Normal Form (2NF)

- **Must satisfy 1NF**

- **No partial dependency**

All non-key attributes must be fully functionally dependent on the entire primary key.

- ❌ `Course_Order(course_id, order_id, course_name)`
(Here, `course_name` depends only on `course_id`)
- ✅ Move `course_name` to `Course(course_id, course_name)`

- Applies **only if** the table has a **composite primary key**.

Third Normal Form (3NF)

- **Must satisfy 2NF**

- **No transitive dependency**

Non-key attributes should not depend on other non-key attributes.

- ❌ `Account(account_id, username, email)`
(If `username → email`, then `email` is transitively dependent)

- **Every non-key attribute must depend directly and only on the primary key**

Stored Routine

Stored routines

- a block of code that save in database call when want
- benefit: reuse code, security
- **Stored functions**
 - return something
 - call by their name not from CALL, be part of query
- **Stored procedure**
 - don't return value
 - call by CALL
- **Triggers**
 - automaticly run when something happen(insert, update, delete)
 - 2 level (ex.100 row but in 1 statement)
 - Row-level - call 100 times
 - Statement-level - call 1 time
 - create function(that return trigger) first, then create trigger (binding)
 - we can DROP, or ALTER

View - Virtual table

- don't make things faster but, easier to call
- reuse query

- update table -> view updated
- usually be read-only
- high query cost, wanna reduce overhead -> materialize view(use more storage)

EXPLAIN, EXPLAIN ANALYZE

- EXPLAIN shows how a SQL query will be executed (the query plan).
- EXPLAIN ANALYZE actually runs the query and shows the real execution steps and time.

Transaction

- update data on many places -> atomicity
- commit things ._.

Windows Function

- ประมวลผลรายการข้อมูล โดยไม่มีการรวมรายการเหมือน aggregate function

NoSQL

- in this course we will learn Document database which is **MongoDB**
- it is **schema-less** NoSQL document database, Scale well both data volume and network traffic
- terminology
 - table -> collection, row -> document, column -> field, table joins -> \$lookup
- Document schema - JSON object that define structure in document
- pros and cons **Embedding inside document**
 - **Pros**
 - pull everything in one query
 - link data between collection through &lookup -> decrease overhead
 - ensure atomic
 - **Cons**
 - document may be too big -> overhead
 - MongoDB max at 16MB
- pros and cons **Referencing**
 - **Pros**
 - document size not too big
 - less redundancy (not a big deal, if perf ok is ok)
 - **Cons**
 - use more queries or \$lookup

schema design rules

1. use embedding first if possible, and one-to-few case too (one-to-many or call separately better use referencing)
2. use case that need to call separately, make it new collection
3. don't use lookup if not necessary
4. one-to-squillions(1000)(ex.log) -> new document and point back to host

5. many-to-many -> new document, point back and forth

Storage and Indexing

Overview

- how does DBMS store and access persistent data?
- minimizing I/O cost
- hash, tree based index?

External Storage

- Disks
 - Random access devices
 - retrieve any page at fixed cost
- Tapes
 - Sequential access device

Files and access methods layer

- Operation support (insert, delete, scan)
- keep track of pages allocation
- tracks available space
- how is a relation stored?
 - as a file of records, each has a **record id** which use to locate page number
 - implemented by the component: Files and access methods layer

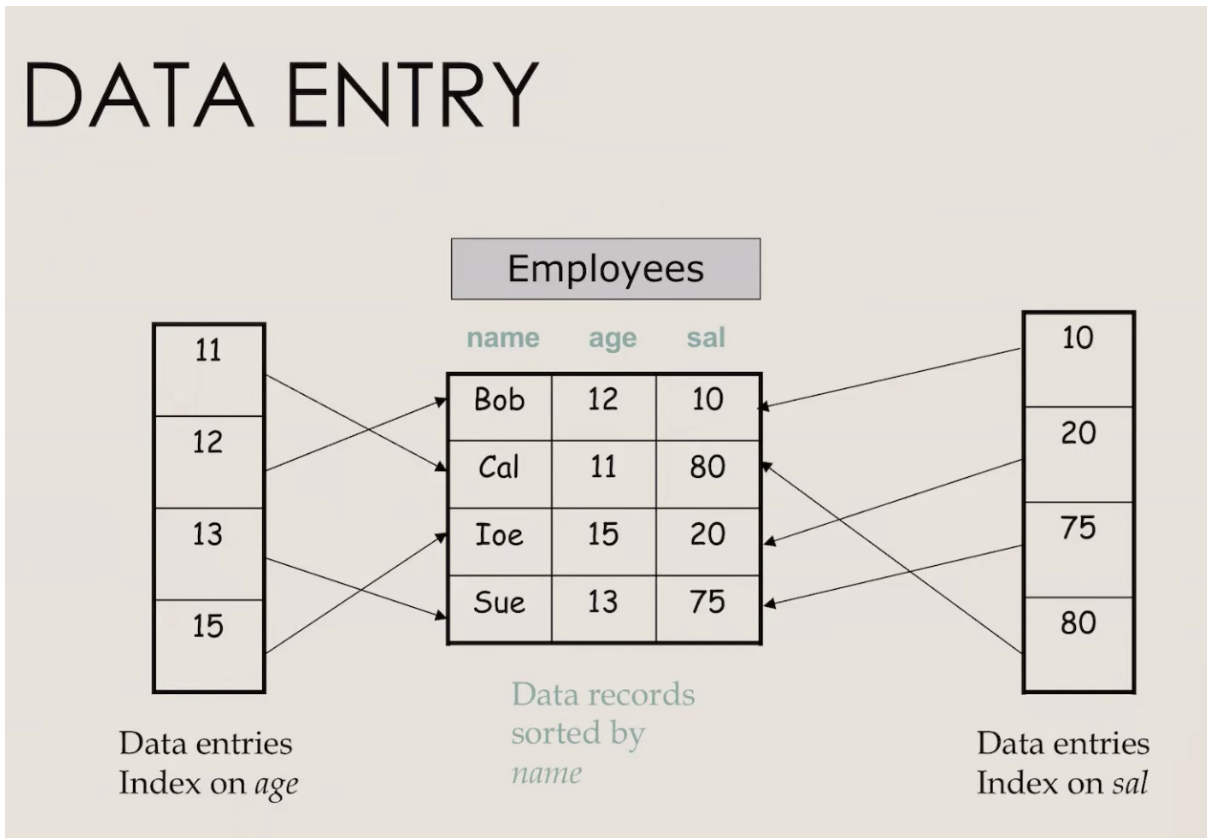
File Organization

- Alternative file organizations
 - heap (random order)
 - sorted
 - indexes

Concepts

- **index file** contains a collection of **data entries** (aka. index entries) -> (search-key, pointer)
- Two basic kinds of indices
 - **Ordered indices**

- Hash indices



Alternative For Data Entries

- **Alternative 1**
 - data entry k^* is an **actual data record**
 - at most 1 index
 - if data is large, # of pages is high
 - indexed file organization
- **Alternative 2**
 - data entry is $\langle k, rid \rangle$ pair
- **Alternative 3**
 - Data entry is $\langle k, rid\text{-list} \rangle$
- 2,3 smaller than 1

Index classification

-
- **Primary index** - search key contains primary key
 - Sequential scan is efficient
 - **Secondary index** - not primary. __.
 - must be dense
 - Sequential scan is expensive
 - (unique index - search key contains candidate key)
 - primary, secondary when file modified, every index must be update
-

- **Clustered index** - order of data record is close to order of data entries
- **Unclustered index** - not clustered __.

- **Dense Index Files** - index record appears for every search key value
 - faster
- **Sparse Index Files**
 - only for **Cluster and Primary** index
 - less space and less maintenance overhead (insert, delete)

Index Data Structure

- **HASH-BASE INDEXES**
 - index is a collection of buckets
 - good for equality selection
- **B+ TREE INDEX**
 - 50-100 fanout -> height really short
 - good for range search
 - always balance in height
 - fill factor - min pointer for a node
 - todo: understand how it balance

Cost Model

In this course we ignore CPU cost, for simplicity (I/O cost dominate), ignore pre-fetching too, use Average-case analysis

- B: The number of data page
- R: Number of data records per page
- D: AVG time to read and write disk page
- **Operations to Compare** - Scan, Equality search, Range selection, Insert selection, Insert a record, Delete a record
- **Compare File Organizations**
 - Heap files(random order; insert at eof)
 - Sorted files, sorted on a search key (no index)
 - Clustered B+ tree file on a search key, Alternative 1
 - Heap file with unclustered B+ tree index Assumptions in our analysis
- Heap files
 - equality selection on key
 - exactly one match
- Sorted files
 - files compact after deletions
- Indexes
 - Alt 2 & 3

- data entry size = 10% size of record
- Hash
 - No overflow bucket
 - 80% page occupancy -> file size = 1.25 data size
- Tree
 - 67% occupancy (typical) -> file size = 1.5 data size

data entry \Rightarrow index
 says this page w/ size of record
 \therefore 1 page can hold more 10 times.
 $\rightarrow 10 R$

$B = \text{page}$, $D = \text{time/page}$, $R = \# \text{ of max data records per page in datafile}$

	(a) Scan	(b) Equality	(c) Range	(d) Insert	(e) Delete
(1) Heap random	BD	best = 1 page worst = B page $0.5BD$	BD	- fetch last page (D) - add. Record (CRD time) - write back (CD) $2D$	- locate ($0.5BD$) - eg search - remove - write back (CD) 1 page $0.5BD + D$
(2) Sorted	BD	$(\log_2 B)D$	- search first item ($D \log B$) - how many match page (scan) $D \log B + (\# \text{ match page})D$	- find pos to insert ($D \log B$) - move other (fetch avg half) ($0.5BD$) - add. data in case (CRD time) - write back (CD) $D \log B + BD$	same, just delete $D \log B + BD$
(3) Clustered B^+ tree Alt 1	- 67% occupancy $\rightarrow 1.5B$ $1.5BD$	- total page $1.5B$ - B^+ tree \rightarrow find out F $D \log_{1.5} B$	- search $D \log_{1.5} B + (\# \text{ match page})D$	- find leaf page ($D \log_{1.5} B$) - add. (CRD time) - write back read page (D) - no need to move leave 67% occup $D \log_{1.5} B + D$	same $D \log_{1.5} B + D$
(4) Unclustered Tree index	- scan leaf ($0.15BD$) - for each data entry, fetch real data (BRD) $BRD + 0.15BD$	- locate page of data entry ($D \log_{1.5} B$) - fetch real data 1 page $D \log_{1.5} B + D$	$D \log_{1.5} B$ + ($\# \text{ match record}$)	- insert new record in leaf 4 fetch page empty (D) - insert (CRD), write back (D) - insert CR data entry - find leaf pos ($D \log_{1.5} B$) - add. new data entry - write back (D) $3D + D \log_{1.5} B$	- locate data entry ($D \log_{1.5} B$) - fetch data record (D) - delete data record - write back (D) - delete data entry - write back (D) $3D + D \log_{1.5} B$
(5) Unclustered Hash index	- scan by index - bucket - scan page in index ($0.15BD$) - for each data in bucket fetch real data (BRD) $0.15BD + BRD$	- locate page/bucket of data entry ($D \log_{1.5} B$) - fetch bucket (D) - fetch data record from data file (scan file) (D) $2D$	- repeat equality search - DBMS don't use index cause scan in data file (BD) BD^* not using index	- insert at end - fetch page (D), insert, write back (D) - locate in chain (scan) - read bucket (D) - add. data entry, write back (D) $4D$	- locate data entry cont - read bucket (D) - fetch data record (D) - delete data record, write back (D) - delete data entry, write back (D) $4D$

Several assumptions underlie these (rough) estimates!

(Tree) 67% occupancy $\rightarrow 1.5B$ page
 data entry size 10% $\rightarrow 0.15B$ page
 (# of leaf page of index)
 $\# \text{ of data entry/page} \rightarrow \frac{2R \times 100}{3} = \frac{20}{3}R \approx 6.7R$
 Fetch every page from every data entry
 in every leaf page: $0.15B \times \frac{20}{3}R \times D = BRD$
 leaf page \rightarrow page

(Hash) 80% occupancy $\rightarrow 1.25B$ page
 entry size 10% $\rightarrow 0.125B$
 $\# \text{ of data per page} = 0.8R \times \frac{100}{10} = 8R$
 fetch every thing
 BRD

• Understanding the Workload

- more selective, you better index it

• Choice of indexes

- consider the most important queries in turn
- trade-off: queries faster, updates slower and require disk space
- Where clause
 - **Exact match - hash index**
 - **range query - tree index**
 - cluster may help
- Multi-attribute
 - order of attribute is important for range queries
 - can sometimes enable **index-only strategies** - clustering is not important

• Example

```
SELECT E.dno FROM Emp E WHERE E.age > 40
```

```
/*
```

```
range so b+ tree index on
```

```
how selective -> if most people age > 40 not that selective so not worth
```

is the index cluster?

*/

```
SELECT E.dno, COUNT(*) FROM Emp E WHERE E.age>10 GROUP BY E.dno
```

/*

consider the Group by

if many age > 10 -> not selective -> better not create index on age
group by dno, so Cluster E.dno index may be better -> dont even have to
fetch, use only index very fast

*/

```
SELECT E.dno FROM Emp E WHERE E.hobby='Stamps'
```

/*

equality search but hobby not unique

clustering on E.hobby helps!

*/

Indexs with Composite Search Keys

- in goodnote 😭